

DeBERTa-v3 Natural Language Inference Hallucination Guard

Hallucination control is enforced by a dedicated Validator Agent utilizing DeBERTa-v3 NLI model. Every generated statement is paired as a hypothesis against the retrieved premise context. If the entailment score is below the groundedness threshold of 0.70, the system rejects the answer. In case of ungrounded generation, the system returns: 'I don't have enough info in documents'. This guarantees 0% unverified hallucinations in critical financial and legal domains.

Author: Mohammad Subhan Pasha