

Qdrant Vector Database Cluster Configuration

Qdrant vector engine manages high dimensional embeddings generated by OpenAI text-embedding-3-large. The vector space dimension is set to 3072 using cosine similarity metric. HNSW indexing is configured with $m=16$ and $ef_construct=100$ for sub-10ms query latency. Collections are partitioned by tenant namespace with payload indexing on `source_type` and `uri` metadata.

Author: Mohammad Subhan Pasha