

Embedding-Based Semantic Chunking Specification

Fixed-size character chunking leads to split context and loss of semantic coherent boundaries. PASHA-NEURO-RAG employs embedding-based semantic sentence chunking. Document text is split into sentences, and cosine distance is calculated between consecutive embeddings. New chunks are created at semantic distance peaks or when reaching target token count of 350 tokens.

Author: Mohammad Subhan Pasha