

MOHAMMAD SUBHAN PASHA

Principal AI Engineer - Tailored for Vercel (AI Engineer (Remote))

Email: subhanpasha@nexus-hive.ai | Phone: +1 (555) 789-0100 | GitHub: github.com/mdsubhanpasha | Remote

EXECUTIVE SUMMARY

Targeted AI Systems Architect tailored for Vercel. Specialized in FastAPI, Playwright, Agent. Proven track record across 6 Enterprise AI OS platforms: Built 25-agent MNC Sim, RAGAS 0.96 RAG engine, and VOX-AI <300ms voice pipeline.

PROFESSIONAL EXPERIENCE & OS PROJECTS

Lead AI Workforce & Swarm Architect | PASHA AI LABS (2023 - Present)

- [STAR] Built 25-agent Autonomous MNC Operating System utilizing LangGraph multi-agent orchestrator with strict QA critic loops. Integrated with Architected PASHA-OS 25-Agent Autonomous MNC System with LangGraph meeting orchestrator and sub-800ms API latency.
- [STAR] Architected NEURO-RAG self-correcting hybrid RAG engine with Qdrant vector database, Reciprocal Rank Fusion, and DeBERTa-v3 NLI guard attaining 0.96 RAGAS score. Integrated with Built PASHA-NEURO-RAG self-correcting RAG engine with DeBERTa-v3 NLI guard & Qdrant hybrid BM25 dense retrieval hitting 0.96 RAGAS score.
- [STAR] Engineered VOX-AI customer voice support engine delivering <300ms end-to-end streaming latency across STT, GPT-4o tool calling, and ElevenLabs TTS. Integrated with Engineered VOX-AI real-time voice assistant with Deepgram Nova-2 STT & ElevenLabs streaming TTS achieving <300ms latency.

Senior Autonomous Systems Engineer | NEXUS AI SYSTEMS (2021 - 2023)

- [STAR] Designed PASHA-UNIFIED-OS auto-publishing LinkedIn OS with 4-node LangGraph pipeline driving 40% growth in viral engagement. Integrated with Developed PASHA-UNIFIED-OS LinkedIn personal brand engine with Groq Llama 3.3 and 4-layer auto-engagement pipeline.
- [STAR] Implemented AUTO-GROWTH marketing engine with 5 CrewAI agents managing real-time SERP and GA4 analytics auto-optimization. Integrated with Created AUTO-GROWTH Autonomous AI Marketing Agency featuring 5 CrewAI/LangGraph specialized agents.
- [STAR] Created PASHA-GLASS privacy context AR assistant with on-device SQLite face vector gallery and AES encrypted transient cache. Integrated with Designed PASHA-GLASS privacy-first AR assistant with local face detection, AES encryption, and sub-50ms cosine similarity matching.

TECHNICAL SKILLS

AI & Agents: LangGraph, CrewAI, Groq Llama 3.3 70B, Qdrant Vector DB, RAGAS, DeBERTa-v3 NLI, Tavily, Playwright

Backend & Systems: Python, FastAPI, Streamlit, PostgreSQL, Redis, Docker, Kubernetes, Microservices, C++

Observability: Prometheus, Grafana, Loki, LogQL, OpenTelemetry, Pytest

EDUCATION & CERTIFICATIONS

B.S. in Computer Science & Artificial Intelligence - Top Tier Institute